

# Vision Transformers Tutorial

<https://github.com/eemlcommunity/PracticalSessions2021/tree/main/vision>

**Andreas Steiner**

EEML summer school  
July 8th 2021, Budapest (virtually)

 Google Research

# Outline

1. Flax models
  - > Practical part
2. Vision Transformer in Flax
  - > Practical part
3. Exploring pre-trained ViTs
  - > Practical part

# Flax Models

And a little bit about JAX

# Researchers 🐱 JAX

## Sam Schoenholz

I would say that JAX has significantly (~5x-10x) improved my productivity and has enabled **research that would have been extremely difficult, if not impossible, with existing tools.** [...] Allowed me to write a molecular dynamics simulation in numpy [...] A first pass at the code took an afternoon to write.

## David Sussillo

It took me only a weekend to reimplement my LFADS sequential autoencoder algorithm. It's running on GPU, using vmap and jit'd. [...] The ease is amazing. [...] The **JAX model is easy for my neuroscience audience to understand**, to deploy, and to innovate upon.

## Vikas Sindhwani

JAX is very compelling for trajectory optimization using Iterative LQR whose key bottleneck is linearization and solving LQR problems: we need fast gradients, Jacobians and Hessians. The **benchmark** is on small but realistic problem dimensions, where JAX turns out to be **competitive with a third party C/C++ QP solver.**

## Elliot Creager

[Differentially-private] SGD requires clipping the per-example parameter gradients, which is non-trivial to implement efficiently. [...] We are able to reproduce the MNIST results in the TensorFlow reference implementation at a **30X speedup** for the simple convolutional architecture used in the original DPSGD paper.

# JAX is Numpy

```
import numpy as np
x = np.arange(9_000_000,
              dtype=np.float32)
x = x.reshape([3000, 3000])
x
```

```
array([[0.000000e+00, ...],
       ...,
       [..., 8.999999e+06]])
```

```
%%timeit
(x @ x)
```

1 loop, best of 5: 428 ms per loop

```
import jax.numpy as jnp
x = jnp.arange(9_000_000,
               dtype=jnp.float32)
x = x.reshape([3000, 3000])
x
```

```
DeviceArray([[0.000000e+00, ...],
              ...,
              [..., 8.999999e+06]])
```

```
%%timeit
(x @ x).block_until_ready()
```

1 loop, best of 5: 11.9 ms per loop

# grad : Compute gradients

```
def loss_fn(weights, inputs, targets):  
    outputs = model(weights, inputs)  
    return ((targets - outputs)**2).mean()  
  
def update_step(weights, inputs, targets):  
    grads = jax.grad(loss_fn)(  
        weights, inputs, targets)  
    weights -= 0.1 * grads  
    return weights  
  
weights = init()  
for inputs, targets in data:  
    weights = update_step(  
        weights, input, targets)
```

`grad()` is a *function transformation*.

Multiple signatures:


```
loss, grads = jax.value_and_grad(loss_fn)(  
    weights, inputs, targets)
```

```
(loss, logits), grads = jax.value_and_grad(  
    loss_fn, has_aux=True)(  
    weights, inputs, targets)
```

# JAX-Arguments : Pytrees

<https://jax.readthedocs.io/en/latest/pytrees.html>

- Most JAX functions operate over pytrees of `jax.ndarray` (=leafs)
- Useful functions :  
`jax.tree_util.tree_[un]flatten()`  
`jax.tree_[multi]map()`
- Flax extends pytrees to [`@flax.struct.dataclass`](#)



# Neural nets in JAX

- Unlike stateful framework (e.g. PyTorch)  
Separate computation & variables
- Functional solution : provide `init()` & matching `apply()`  
for *every module*
- OMG lots of code 🤯

## Flax to the rescue

- Opinionated `nn.Module` abstraction
- "init" & "apply" possible in a single function



# Simple PyTorch NN

```
import torch.nn as nn
import torch.nn.functional as F
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv = nn.Conv2d(1, 16, 3, 3)
        self.fc = nn.Linear(1296, 10)
    def forward(self, x):
        x = self.conv(x)
        x = F.relu(x)
        x = torch.flatten(x, 1)
        x = self.fc(x)
        return F.log_softmax(x, dim=1)
```

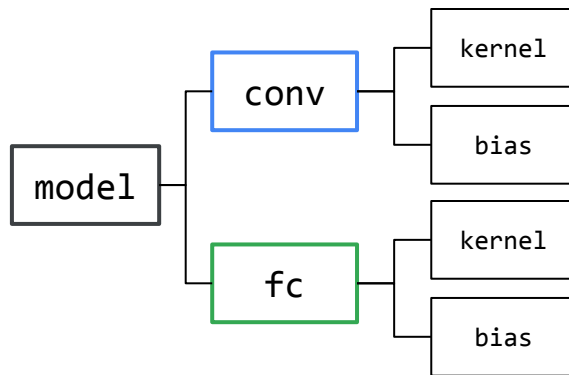
```
model = Net()
optimizer = optim.Adam(model.parameters(), lr=0.1)
optimizer.zero_grad()
outputs = model(inputs)
loss = F.nll_loss(outputs, labels)
loss.backward()
optimizer.step()
```

- Model: setup/forward phase with shape-dependent values
- Dataflow is not clear for the uninitiated (model↔loss↔optimizer)

# Flax "Linen" modules

```
import flax.linen as nn
class Net(nn.Module):
    units: int
    def setup(self):
        self.conv = nn.Conv(
            self.units, [3, 3], [3, 3])
        self.fc = nn.Dense(10)
    def __call__(self, x, *, train):
        x = self.conv(x)
        x = nn.relu(x)
        x = x.reshape([len(x), -1])
        x = self.fc(x)
        return nn.log_softmax(x)
```

- *Immutable* dataclasses.
- Predefined modules like `nn.Conv()`
- Auto shape inference.
- Automatically creates pytree(s) that reflects module hierarchy:



# Linen modules - compact 🕶️

Avoids replicating logic in `__call__()` and `setup()`.

```
class Net(nn.Module):  
    units: int  
    @nn.compact  
    def __call__(self, x):  
        x = nn.Conv(self.units, [3, 3], [3, 3])(x)  
        x = nn.relu(x)  
        x = x.reshape([len(x), -1])  
        x = nn.Dense(10)(x)  
        return nn.log_softmax(x)
```

# Using Linen modules

```
model = Model(num_classes=num_classes)
variables = model.init(key, inputs)
```

1. Instantiate module
2. Create initial variables

```
outputs = model.apply(variables, inputs)
```

No mutable state:

$$f(v_{in}, x) \rightarrow y$$

# Practical Part

[https://github.com/eemlcommunity/  
PracticalSessions2021/tree/main/vision](https://github.com/eemlcommunity/PracticalSessions2021/tree/main/vision)

(Load data)

Define a model

Create and use model weights

Train, evaluate + predict

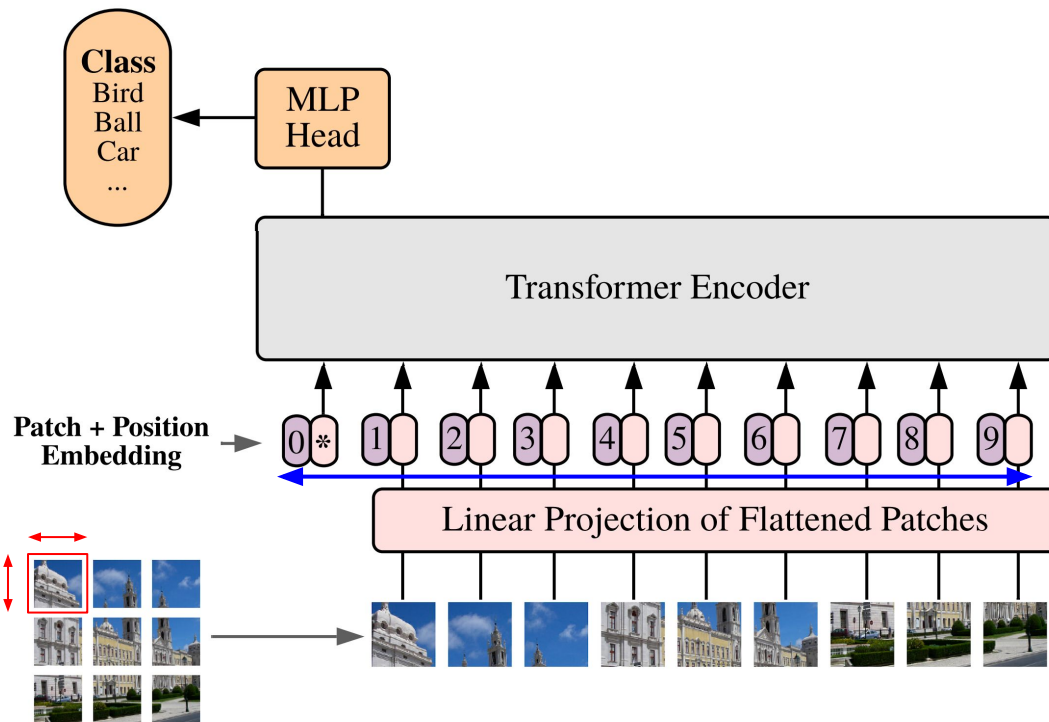
# Vision Transformers in Flax

# Vision Transformer (ViT)

**Idea:** Take a transformer and apply it directly to image patches

sequence\_length  
depends on image size

patch\_size  
e.g.  $B/32$

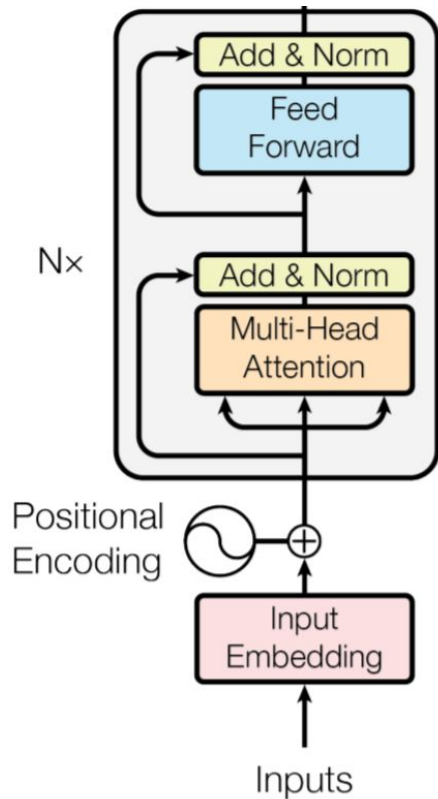


[Cordonnier et al., On the Relationship between Self-Attention and Convolutional Layers, ICLR 2020](#)

[Dosovitskiy et al., An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale, ICLR 2021](#)

# Transformer


- Transformer “encoder”
  - A stack of alternating self-attention and MLP blocks
  - Residuals and LayerNorm
- Transformer “decoder” (not shown)
  - A slightly more involved architecture useful when the output space is different from the input space (e.g. translation)





# Self-attention

- Each of the tokens (=vectors) attends to all tokens
  - Extra tricks: learned key, query, and value projections, inverse-sqrt scaling in the softmax, and multi-headed attention (omit for simplicity)
- It's a set operation (permutation-invariant)
  - ...and hence need “position embeddings” to “remember” the spatial structure
- It's a global operation
  - Aggregates information from all tokens



$$\alpha_j = \text{softmax}\left(\frac{K \mathbf{x}_1 \cdot Q \mathbf{x}_j}{\sqrt{d_K}}, \dots, \frac{K \mathbf{x}_n \cdot Q \mathbf{x}_j}{\sqrt{d_K}}\right)$$
$$\mathbf{y}_j = \sum_{i=1}^n \alpha_{ji} V \mathbf{x}_i$$

**Simplified!** Multi-headed attention not shown

# Scaling with Data

## Key


ViT = Vision Transformer

BiT = Big Transfer (~ResNet)

ViT overfits on ImageNet, but shines on larger datasets.

\* with heavy regularization ViT has been shown to also work on ImageNet (Touvron et al.)

\*\* training ViT on ImageNet with the sharpness-aware minimizer (SAM) also works very well (Chen et al.)



[Dosovitskiy et al., An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale, ICLR 2021](#)

[Xiangning Chen et al., When Vision Transformers Outperform ResNets without Pretraining or Strong Data Augmentations, arXiv 2021](#)

[Touvron et al., Training data-efficient image transformers & distillation through attention, arXiv 2020](#)

# Practical Part

[https://github.com/eemlcommunity/  
PracticalSessions2021/tree/main/vision](https://github.com/eemlcommunity/PracticalSessions2021/tree/main/vision)

Forming patches + position embeddings

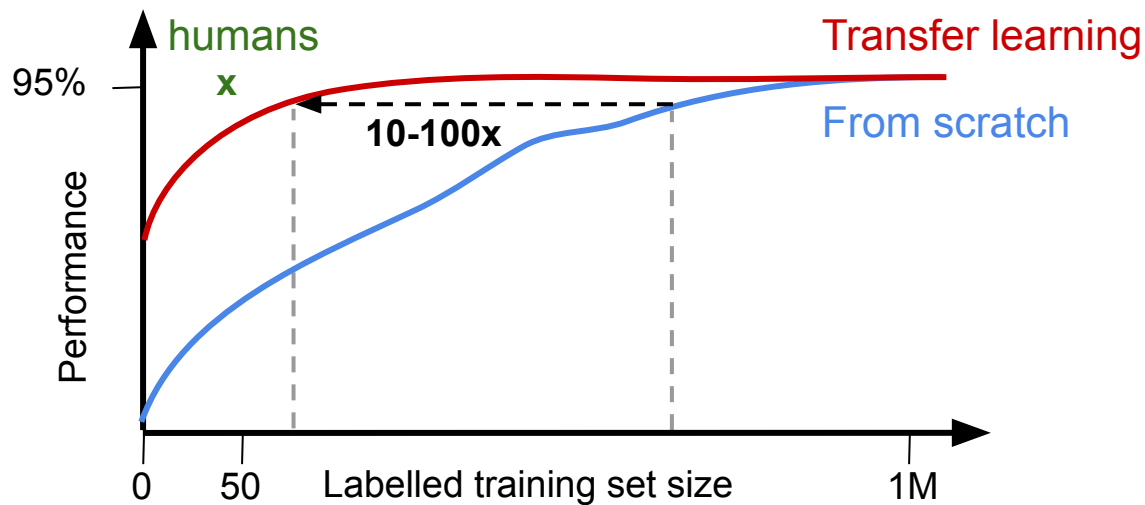
TransformerLayer + Transformer

Understanding ViT params

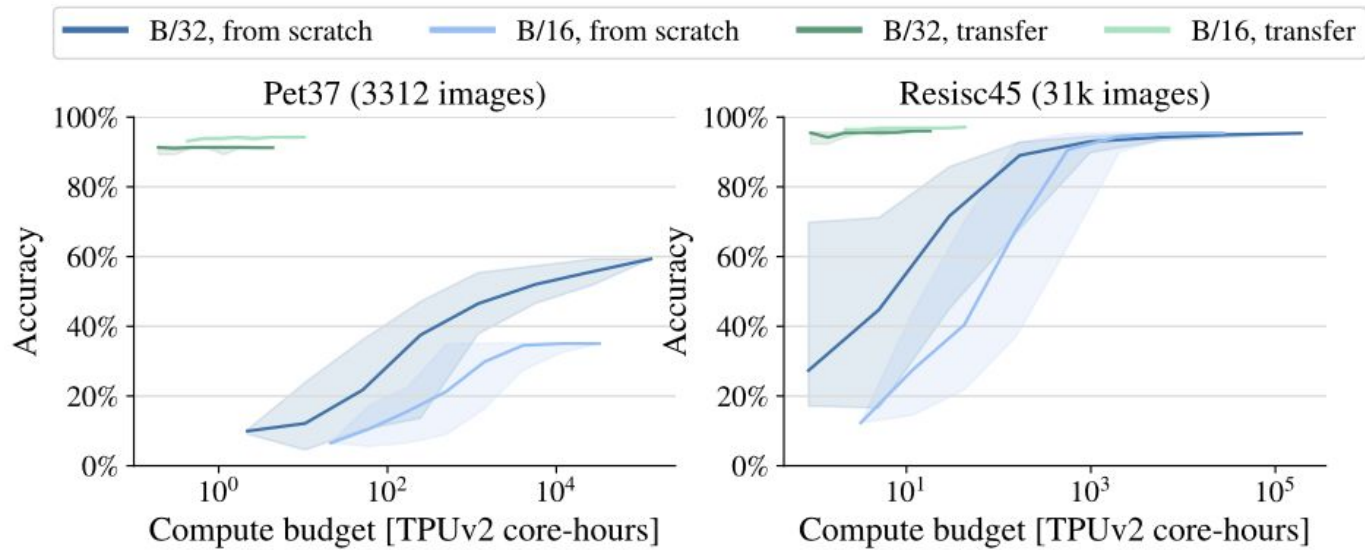
Refined training

# Exploring pre-trained ViTs

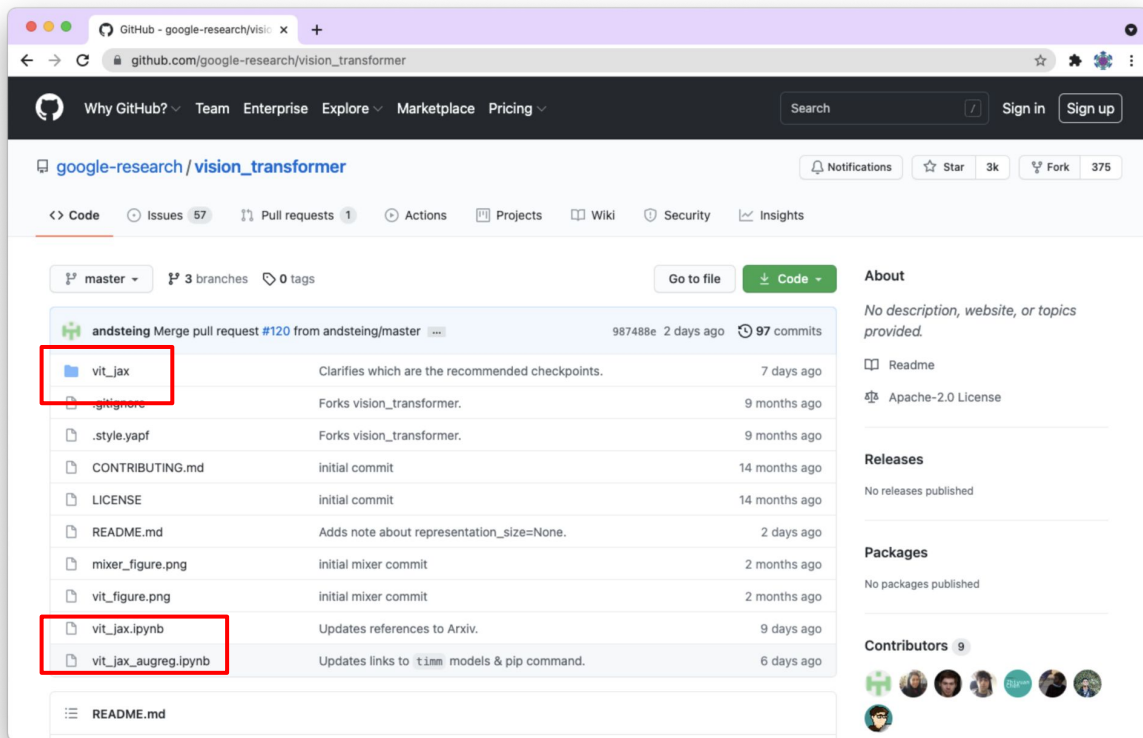
# Why transfer?




# Why transfer?



# vision\_transformer Github repo



# >50k Pre-trained checkpoints



```
# Upstream AugReg parameters (section 3.3):
(
df.groupby(['ds', 'name', 'wd', 'do', 'sd', 'aug']).filename
  .count().unstack().unstack().unstack()
  .dropna(1, 'all').astype(int)
  .iloc[:7] # Just show beginning of a long table.
)
```

			aug	light0		light1		medium1		medium2		none		strong1		strong2	
sd			0.0	0.1	0.0	0.1	0.0	0.1	0.0	0.1	0.0	0.1	0.0	0.1	0.0	0.1	0.1
do			0.0	0.1	0.0	0.1	0.0	0.1	0.0	0.1	0.0	0.1	0.0	0.1	0.0	0.1	0.1
ds	name	wd															
i1k	B/16	0.03	68	68	68	68	68	68	68	68	68	68	68	68	68	68	68
		0.10	68	68	68	68	68	68	68	68	68	68	68	68	68	68	68
	B/32	0.03	68	68	68	68	68	68	68	68	68	68	68	68	68	68	68
		0.10	68	68	68	68	68	68	68	68	68	68	68	68	68	68	68
	L/16	0.03	68	68	68	68	68	68	68	68	68	68	68	68	68	68	68
		0.10	68	68	68	68	68	68	68	68	68	68	68	68	68	68	68
	R+Ti/16	0.03	68	68	68	68	68	68	68	68	68	68	68	68	68	68	68



# Practical Part

[https://github.com/eemlcommunity/  
PracticalSessions2021/tree/main/vision](https://github.com/eemlcommunity/PracticalSessions2021/tree/main/vision)

Use repository code in Colab

Explore checkpoints

Load checkpoints + inference

Fine-tune checkpoints